

Ex. No	For a given set of training data examples stored in a .CSV file, implement and demonstrate the Candidate-Elimination algorithm to output a description of the set of all hypotheses consistent with the training examples	Date
1		

Aim:

The aim is to implement the Candidate-Elimination algorithm to find all hypotheses that are consistent with a given set of training data examples stored in a CSV file. This involves maintaining and updating the sets of the most specific (S) and the most general (G) hypotheses throughout the learning process.

Algorithm:

Initialization: Start with the most specific hypothesis (S) and the most general hypothesis (G).(S) is initialized to the most specific hypothesis (all elements are the null value).(G) is initialized to the most general hypothesis (all elements are the wildcard '?').

Iterate Through Training Examples: For each positive example, update (S) to be more general to cover the example, and remove inconsistent hypotheses from (G). For each negative example, update (G) to be more specific to exclude the example, and remove inconsistent hypotheses from (S).

Consistent Hypotheses: At the end of the process, (S) and (G) represent the set of hypotheses consistent with the training examples.

Program:

```

import numpy as npimport pandas as pd
data=pd.DataFrame(data=pd.read_csv('enjoysport.csv'))concepts=np.array(data.iloc[:,0:-1])
print(concepts)
target = np.array(data.iloc[:, -1])print(target)

def learn(concepts, target):specific_h=concepts[0].copy()
    print("initialization of specific_h and general_h")print(specific_h)
    general_h=[["?" for i in range(len(specific_h))] for i in range(len(specific_h))]
    print(general_h)
    for i, h in enumerate(concepts):if target[i] == "yes":
        for x in range(len(specific_h)):if h[x] != specific_h[x]:
    
```

```

specific_h[x] = '?' general_h[x][x] = '?'
print(specific_h) print(specific_h)

if target[i] == "no":
    for x in range(len(specific_h)):
        if h[x] != specific_h[x]:
            general_h[x][x] = specific_h[x]
        else:
            general_h[x][x] = '?'
    print(" " + str(i) + " steps of Candidate Elimination Algorithm", i+1)
    print(general_h)

indices = [i for i, val in enumerate(general_h) if val == ['?', '?', '?', '?', '?', '?']]
for i in indices:
    general_h.remove(['?', '?', '?', '?', '?', '?'])
return specific_h, general_h

s_final, g_final = learn(concepts, target)
print("Final Specific_h: ", s_final, sep="\n")
print("Final General_h: ", g_final, sep="\n")

```

Data Set:

Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
sunny	warm	normal	strong	warm	same	yes
sunny	warm	high	strong	warm	same	yes
rainy	cold	high	strong	warm	change	no
sunny	warm	high	strong	cool	change	yes

Output:

FinalSpecific_h:
['sunny' 'warm' '?' 'strong' '?' '?']

FinalGeneral_h:

[['sunny', '?', '?', '?', '?', '?'],
['?', 'warm', '?', '?', '?', '?']]

Result:

Thus the program was successfully executed.

Ex. No	Write a program to demonstrate the working of the decision tree based ID3 algorithm. Use an appropriate data set for building the decision tree and apply this knowledge to classify a new sample.	Date
2		

Aim:

To implement and demonstrate the working of the decision tree-based ID3 (Iterative Dichotomiser 3) algorithm. This involves building a decision tree from a given dataset and using the constructed tree to classify new samples.

Algorithm:

Select Attribute: Choose the best attribute (A) to split the dataset based on a criterion (e.g., information gain or Gini impurity).

Split Dataset: Split the dataset into subsets based on the values of attribute (A).

Recursive Build: Recursively apply the above steps to each subset until one of the following conditions is met: All instances in the subset belong to the same class. There are no more attributes to split on. Stopping criteria are met (e.g., maximum tree depth reached).

Create Decision Tree: Create a decision tree node that represents the selected attribute (A).

TrainingDataset:

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

TestDataset:

Day	Outlook	Temperature	Humidity	Wind
T1	Rain	Cool	Normal	Strong
T2	Sunny	Mild	Normal	Strong

Program:

```
import mathimportcsv

defload_csv(filename):

    nes=csv.reader(open(filename,"r"));
    dataset = list(lines)headers = dataset.pop(0)returndataset,headers

classNode:

    definit(self,attribute):self.attribute=attribute self.children=[] self.answer=""

defsubtables(data,col,delete):

    c={}
    coldata=[row[col] for row in data]attr=list(set(coldata))

    counts=[0]*len(attr)r=len(data)c=len(data[0])
    for x in range(len(attr)):for y in range(r):
        if data[y][col]==attr[x]:counts[x]+=1

    for x in range(len(attr)):
        dic[attr[x]]=[[0 for i in range(c)] for j in range(counts[x])]

    pos=0
    for y in range(r):
        if data[y][col]==attr[x]:if delete:
            del data[y][col]dic[attr[x]][pos]=data[y]pos+=1
```

```
returnattr,dic
```

```
defentropy(S):
```

```
tr=list(set(S))
```

```
if len(attr)==1: return0
```

```
counts=[0,0]
```

```
for i in range(2):
```

```
counts[i]=sum([1 for x in S if attr[i]==x])/(len(S)*1.0)
```

```
sums=0
```

```
for cnt in counts:
```

```
sums+=-1*cnt*math.log(cnt,2) return sums
```

```
defcompute_gain(data,col):
```

```
tr,dic=subtables(data,col,delete=False)
```

```
total_size=len(data) entropies=[0]*len(attr) ratio=[0]*len(attr)
```

```
total_entropy=entropy([row[-1] for row in data]) for x in range(len(attr)):
```

```
ratio[x]=len(dic[attr[x]])/(total_size*1.0) entropies[x]=entropy([row[-1] for row in dic[attr[x]]])
```

```
total_entropy-=ratio[x]*entropies[x] return total_entropy
```

```
defbuild_tree(data,features):
```

```
stcol=[row[-1] for row in data] if len(set(lastcol))==1:
```

```
node=Node("") node.answer=lastcol[0] return node
```

```
n=len(data[0])-1 gains=[0]*n
```

```
for col in range(n): gains[col]=compute_gain(data,col)
```

```

split=gains.index(max(gains))node=Node(features[split])
fea = features[:split]+features[split+1:]attr,dic=subtables(data,split,delete=True)

for x in range(len(attr)):child=build_tree(dic[attr[x]],fea)node.children.append((attr[x],child))
returnnode

defprint_tree(node,level):
    if node.answer!="":
        print("    "*level,node.answer)return

    print("    "*level,node.attribute)forvalue,ninnode.children:
        print("        "*(level+1),value)print_tree(n,level+2)

defclassify(node,x_test,features):
    ifnode.answer!="":
        print(node.answer)return

    pos=features.index(node.attribute)forvalue,ninnode.children:
        if x_test[pos]==value:classify(n,x_test,features)

"Mainprogram"
dataset,features=load_csv("data3.csv")node1=build_tree(dataset,features)

print("The decision tree for the dataset using ID3 algorithmis")
print_tree(node1,0)testdata,features=load_csv("data3_test.csv")forxtest intestdata:
    print("Thetestinstance:",xtest)
    print("Thelabelfortestinstance:",end=" ")classify(node1,xtest,features)

```

Output:

The decision tree for the dataset using ID3 algorithmis

rain
wind
strong
weak
overcast
yes

sunny
normal
high

The test instance: ['rain', 'cool', 'normal', 'strong']

The label for test instance: no

The test instance: ['sunny', 'mild', 'normal', 'strong']

The label for test instance: yes



Result:

Thus the program was successfully executed.

Ex. No	Build an Artificial Neural Network by implementing the Back propagation algorithm and test the same using appropriate data sets.	Date
3		

Aim:

To implement an Artificial Neural Network (ANN) using the Backpropagation algorithm and test its performance using appropriate datasets. This involves training the neural network on a dataset to learn patterns and then using the trained network to make predictions or classifications on new data.

Algorithm:

Load Dataset: Load a dataset suitable for training and testing the ANN. The dataset should have input features and corresponding target labels.

Initialize Network: Initialize the neural network with random weights. **Forward Pass:** Implement the forward pass to compute the output of the network given input data.

Compute Loss: Compute the loss or error between the predicted output and the actual output using a suitable loss function (e.g., mean squared error).

Backward Pass: Implement the backward pass to compute the gradients of the loss function with respect to the weights.

Update Weights: Use gradient descent or its variants to update the weights of the network based on the computed gradients.

Training Loop: Iteratively perform forward and backward passes on the training data for a fixed number of epochs.

Test the ANN: After training, evaluate the performance of the ANN on a separate test dataset to assess its generalization ability.

Training Examples:

Example	Sleep	Study	Expected %in Exams
1	2	9	92
2	1	5	86
3	3	6	89

Normalizetheinput

Example	Sleep	Study	Expected %inExams
1	$2/3=0.66666667$	$9/9=1$	0.92
2	$1/3=0.33333333$	$5/9=0.55555556$	0.86
3	$3/3=1$	$6/9=0.66666667$	0.89

Program:

```
import numpy as np
X=np.array(([2,9],[1,5],[3,6]),dtype=float)
y=np.array(([92],[86],[89]),dtype=float)
X = X/np.amax(X, axis=0) # maximum of X array longitudinally
y = y/100
```

```
#Sigmoid Function
def sigmoid(x):
```

```
    return 1/(1+np.exp(-x))
```

```
#Derivative of Sigmoid Function
def derivatives_sigmoid(x):
```

```
    return x*(1-x)
```

```
#Variable initialization
```

```
epoch=5000 #Setting training iterations
lr=0.1 #Setting learning rate
inputlayer_neurons=2 #number of features in data
hiddenlayer_neurons=3 #number of hidden layers neurons
output_neurons=1 #number of neurons at output layer
```

```
#weight and bias initialization
wh=np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))
bh=np.random.uniform(size=(1,hiddenlayer_neurons))
wout=np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
```

```
bout=np.random.uniform(size=(1,output_neurons))
```

```
#draws a random range of numbers uniformly of dim x*y for i in range(epoch):
```

```
#Forward Propagation
hinp1=np.dot(X,wh)
hinp=hinp1+bh
```

```
hlayer_act = sigmoid(hinp)
outinp1=np.dot(hlayer_act,wout)
outinp=outinp1+bout
```

```

output=sigmoid(outinp)

#BackpropagationEO=y-output
    outgrad = derivatives_sigmoid(output)d_output=EO* outgrad
    EH=d_output.dot(wout.T)

#how much hidden layer wts contributed to errorhiddengrad =
    derivatives_sigmoid(hlayer_act)d_hiddenlayer=EH *hiddengrad

# dotproduct of nextlayererror and currentlayeropwout+=hlayer_act.T.dot(d_output)*lr
    wh+=X.T.dot(d_hiddenlayer)*lr

print("Input: \n" + str(X))print("Actual Output: \n" + str(y))print("PredictedOutput:\n",output)

```

Output:

Input:

```

[[0.666666671.          ]
 [0.33333330.55555556]
 [1.          0.66666667]]

```

Actual Output: [[0.92]

```

[0.86]
[0.89]]

```

Predicted Output: [[0.89726759]

```

[0.87196896]
[0.9000671]]

```

Result:

Thus the program was successfully executed.

Ex. No	Write a program to implement the naïve Bayesian classifier for a sample training data set stored as a .CSV file and compute the accuracy with a few test data sets.	Date
4		

Aim:

To implement a Naive Bayes classifier to classify data based on a sample training dataset stored in a CSV file. The Naive Bayes classifier is a simple probabilistic classifier based on Bayes' theorem with the assumption of independence between every pair of features.

Algorithm:

Load Dataset: Load the dataset from a CSV file.

Calculate Prior Probabilities: Calculate the prior probabilities for each class based on the training data.

Calculate Likelihoods: For each feature, calculate the likelihood given each class.

Compute Posterior Probabilities: Use the likelihoods and priors to compute the posterior probabilities.

Predict Class: Classify new samples based on the computed posterior probabilities.

Program:

```

import csvimport randomimportmath
defloadcsv(filename):
    lines = csv.reader(open(filename, "r"));dataset= list(lines)
    foriinrange(len(dataset)):
        #convertingstringstointegersforprocessing
    returndataset
defsplitdataset(dataset,splitratio):
    #67%trainingsize
    trainset=[]
    copy=list(dataset);

```

```
#generateindicesforthedatasetlistrandomlytopickelefortrainingdata
while len(trainset) < trainsize:
    index = random.randrange(len(copy)); trainset.append(copy.pop(index))
return [trainset, copy]
```

```
def separate_by_class(dataset):
```

```
#creates a dictionary of classes 1 and 0 where the values are #the instances belonging to each class
separated = {} #dictionary of classes 1 and 0
for i in range(len(dataset)):
    vector = dataset[i]
    if (vector[-1] not in separated):
        separated[vector[-1]] = []
    separated[vector[-1]].append(vector)
return separated
```

```
def mean(numbers):
```

```
    return sum(numbers) / float(len(numbers))
```

```
def stdev(numbers):
```

```
    avg = mean(numbers)
    variance = sum([pow(x - avg, 2) for x in numbers]) / float(len(numbers) - 1)
    return math.sqrt(variance)
```

```
def summarize(dataset): #creates a dictionary of class summaries
```

```
    [(mean(attribute), stdev(attribute))
     for
```

```
     attribute in zip(*dataset)];
```

```
    del summaries[-1] #excluding labels +ve or -ve
    return summaries
```

```
def summarize_by_class(dataset):
```

```
    separated = separate_by_class(dataset); #print(separated)
```

```
    summaries = {}
```

```
    for class_value, instances in separated.items():
```

```
#for key, value in dic.items():
```

```
#summaries is a dict of tuples (mean, std) for each class value
```

```

        summaries[classvalue] = summarize(instances)#summarize is used to calculate mean and standard deviation
to mean and std

    return summaries

def calculateprobability(x, mean, stdev): exponent=math.exp(-(math.pow(x-mean,2)/
(2*math.pow(stdev,2))))
    return(1/(math.sqrt(2*math.pi)*stdev))*exponent

def calculateclassprobabilities(summaries,inputvector):
    #probabilities contain the all probability of all classes of test data
    probabilities = {}#probabilities contain the all probability of all classes of test data
    for classvalue, classsummaries in summaries.items():#class and attribute information as mean and
sd
        probabilities[classvalue]=1
        for i in range(len(classsummaries)):
            mean, stdev = classsummaries[i] #take mean and standard deviation for every attribute for class 0 and 1 separately
and standard deviation for every attribute for class 0 and 1 separately
            x = inputvector[i] #test vector's first attribute
            probabilities[classvalue]*= calculateprobability(x, mean, stdev);#use normal distribution
            return probabilities

def predict(summaries, inputvector): #training and test data is passed
    probabilities = calculateclassprobabilities(summaries, inputvector)
    bestLabel, bestProb = None, -1
    for classvalue, probability in probabilities.items():#assigns that class which has the highest probability
        if bestLabel is None or probability > bestProb: bestProb = probability
        bestLabel = classvalue
    return bestLabel

def getpredictions(summaries, testset): predictions = []
    for i in range(len(testset)):
        result = predict(summaries, testset[i])
        predictions.append(result)
    return predictions

def getaccuracy(testset, predictions): correct = 0

```

```
for i in range(len(testset)):  
    if testset[i][-1] == predictions[i]: correct+= 1  
return(correct/float(len(testset)))*100.0  
  
def main():  
    filename = 'naivedata.csv' splitratio = 0.67  
    dataset = loadcsv(filename);  
  
    trainingset, testset = splitdataset(dataset,  
    splitratio) print('Split {0} rows into train = {1} and test = {2}'  
    .format(len(dataset), len(trainingset), len(testset))) # prepare model  
    summaries = summarizebyclass(trainingset); # print(summaries)  
    # test model  
    predictions = getpredictions(summaries, testset) # find  
    the predictions of test data with the training data  
    accuracy = getaccuracy(testset, predictions) print('Accuracy of the classifier is:  
    {0}'.format(accuracy)) main()
```

Output:

```
Split 768 rows into train = 514  
and test = 254 rows Accuracy of the classifier is: 71.65354330708661%
```

Result:

Thus the program was successfully executed.

Ex. No	Implement naïve Bayesian Classifier model to classify a set of documents and measure the accuracy, precision, and recall.	Date
5		

Aim:

To implement a Naive Bayes Classifier to classify a set of documents and measure its performance using metrics such as accuracy, precision, and recall. Naive Bayes is a probabilistic classifier based on Bayes' theorem with the assumption of independence between features.

Algorithm:

Preprocessing: Convert the documents into a suitable format (e.g., bag-of-words model). Split the dataset into training and testing sets.

Calculate Prior Probabilities: Compute the prior probability of each class based on the training data.

Calculate Likelihood: For each feature, calculate the likelihood of that feature given a class.

Apply Bayes' Theorem: Use Bayes' theorem to compute the posterior probability of each class given a document.

Classify: Assign the class with the highest posterior probability to the document.

Evaluate Performance: Measure the accuracy, precision, and recall of the classifier using the test data.

Dataset:

	TextDocuments	Label
1	Ilovethissandwich	pos
2	Thisisanamazing place	pos
3	Ifeelverygoodabouttheseebeers	pos
4	Thisismy bestwork	pos
5	Whatanawesomeview	pos
6	Idonotlikethisrestaurant	neg
7	Iamtiredofthisstuff	neg
8	Ican'tdealwiththis	neg
9	Heismyswornenemy	neg

10	Mybossishorrible	neg
11	Thisisanawesome place	pos
12	I donotlikethetasteofthis juice	neg
13	I love todance	pos
14	Iamsickandtired ofthisplace	neg
15	Whatagreatholiday	pos
16	That is a bad locality to stay	neg
17	We will have good fun tomorrow	pos
18	I went to my enemy's house today	neg

Program:

```

import pandas as pd
msg=pd.read_csv('naivetext.csv',names=['message','label'])
print('The dimensions of the dataset',msg.shape)

msg['labelnum']=msg.label.map({'pos':1,'neg':0})
X=msg.messagey=msg.labelnum

print(X)print(y)

#splitting the dataset into train and test data
from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(X,y)

print("\n The total number of Training Data
:",ytrain.shape)print("\n The total number of Test Data:",ytest.shape)

#output of count vectoriser is a sparse matrix
from sklearn.feature_extraction.text import CountVectorizer
count=CountVectorizer()
xtrain_dtm=count.fit_transform(xtrain)xtest_dtm=count_vect.transform(xtest)
print("\n The words or Tokens in the text documents \n")print(count_vect.get_feature_names())

```

```
df=pd.DataFrame(xtrain_dtm.toarray(),columns=count_vect.get_feature_names())

#TrainingNaiveBayes(NB)classifierontrainingdata.
from sklearn.naive_bayes import MultinomialNB
clf = MultinomialNB().fit(xtrain_dtm,ytrain)predicted=clf.predict(xtest_dtm)

#printingaccuracy,Confusionmatrix,PrecisionandRecall
from sklearn import metrics
print('\nAccuracy of the classifier is',
metrics.accuracy_score(ytest,predicted))

print("\n Confusion matrix")print(metrics.confusion_matrix(ytest,predicted))
print("\n The value of Precision' ,metrics.precision_score(ytest,predicted))
print("\n The value of Recall' ,metrics.recall_score(ytest,predicted))
```

Output:

The dimensions of the dataset(18,2)

```
0    Ilovethissandwich
1    Thisisanamazing place
2    Ifeelverygoodaboutthesebeers
3    Thisismy bestwork
4    Whatanawesomeview
5    Idonotlikethisrestaurant
6    Iamtiredofthisstuff
7    Ican'tdealwiththis
8    Heismyswornenemy
9    Mybossishorrible
10   Thisisanawesome place
11   Idonotlikethe tasteofthisjuice
12   Ilove todance
13   Iamsickandtiredofthisplace
14   Whatagreatholiday
15   Thatisabadlocalityto stay
16   Wewillhavegood fun tomorrow
17   Iwenttomyenemy'shousetoday
```

Name: message, dtype: object0 1

1 1
2 1
3 1
4 1
5 0
6 0
7 0
8 0
9 0
10 1
11 0
12 1
13 0
14 1
15 0
16 1
17 0

Name:labelnum,dtype:int64

The total number of Training Data: (13,) The total number of TestData: (5,)

The words or tokens in the text documents

['about', 'am', 'amazing', 'an', 'and', 'awesome', 'beers', 'best', 'can', 'deal', 'do', 'enemy', 'feel',
'fun', 'good', 'great', 'have', 'he', 'holiday', 'house', 'is', 'like', 'love', 'my', 'not', 'of', 'place',
'restaurant', 'sandwich', 'sick', 'sworn', 'these', 'this', 'tired', 'to', 'today', 'tomorrow', 'very', 'view', 'we', 'went',
'what', 'will', 'with', 'work']

Accuracy of the classifier is 0.8Confusionmatrix

[[21]

[02]]

The value of Precision 0.6666666666666666 The value of Recall 1.0

Result:

Thus the program was successfully executed.

Ex. No	Write a program to construct a Bayesian network to diagnose CORONA infection using standard WHO Data Set.	Date
6		

Aim:

To construct a Bayesian Network to diagnose COVID-19 infection using a standard dataset. A Bayesian Network is a graphical model that represents the probabilistic relationships among a set of variables. In the context of COVID-19 diagnosis, it can be used to model the relationships between various symptoms and the likelihood of infection.

Algorithm:

Define Variables: Identify and define the variables involved, such as symptoms (e.g., fever, cough, fatigue), test results, and the infection status.

Structure the Network: Define the structure of the Bayesian Network by determining which variables are conditionally dependent on each other.

Parameter Learning: Estimate the conditional probability tables for each variable given its parents in the network. This can be done using a dataset.

Inference: Use the Bayesian Network to perform inference, i.e., compute the probability of infection given observed symptoms and test results.

Program:

```

import numpy as npimport pandas as pdimportcsv
frompgmpy.estimatorsimportMaximumLikelihoodEstimatorfrompgmpy.modelsimportBa
yesianModel
frompgmpy.inferenceimportVariableElimination

#readClevelandHeartDiseasedata
heartDisease=pd.read_csv('heart.csv')
heartDisease=heartDisease.replace('?',np.nan)

#displaythedata
print('Sample instances from the dataset are given below')print(heartDisease.head())

#displaytheAttributenamesanddatatypes
print("\n Attributes and datatypes")print(heartDisease.dtypes)

```

```
#CreateModel-BayesianNetwork
model = BayesianModel([('age','heartdisease'),('sex','heartdisease'),
('exang','heartdisease'),('cp','heartdisease'),('heartdisease','restecg'),
('heartdisease','chol')])
```

```
#LearningCPDsusingMaximumLikelihoodEstimators
```

```
print("\n Learning CPD using Maximum likelihood
estimators")model.fit(heartDisease,estimator=MaximumLikelihoodEstimator)
```

```
#InferencingwithBayesianNetwork
```

```
print("\n Inferencing with Bayesian
Network:")HeartDiseasetest_infer=VariableElimination(model)
```

```
#computingtheProbabilityofHeartDiseasegivenrestecg
```

```
print("\n1. Probability of HeartDisease given evidence = restecg
:1")q1=HeartDiseasetest_infer.query(variables=['heartdisease'],evidence={'restecg':1})
print(q1)
```

```
#computingtheProbabilityofHeartDiseasegivencp
```

```
print("\n2. Probability of HeartDisease given evidence = cp:2")q2=HeartDiseasetest_infer.query(vari
ables=['heartdisease'],evidence={'cp':2})
print(q2)
```

Output:

```
===== RESTART: E:\ML Lab - 2020-21\MLLab-1\ML1.py =====
```

```
Few examples from the dataset are given below
```

	age	sex	cp	trestbps	chol	...	oldpeak	slope	ca	thal	heartdisease	
0	63	1	1	145	233	...	2.3	3	0	6		0
1	67	1	4	160	286	...	1.5	2	3	3		2
2	67	1	4	120	229	...	2.6	2	2	7		1
3	37	1	3	130	250	...	3.5	3	0	3		0
4	41	0	2	130	204	...	1.4	1	0	3		0

```
[5 rows x 14 columns]
```

```
Attributes and datatypes
```

age	int64
sex	int64
cp	int64
trestbps	int64
chol	int64
fbs	int64
restecg	int64
thalach	int64
exang	int64
oldpeak	float64
slope	int64
ca	object
thal	object
heartdisease	int64
dtype:	object

Learning CPD using Maximum likelihood estimators

Inferencing with Bayesian Network:

1. Probability of HeartDisease given evidence= restecg

heartdisease	phi(heartdisease)
heartdisease(0)	0.1012
heartdisease(1)	0.0000
heartdisease(2)	0.2392
heartdisease(3)	0.2015
heartdisease(4)	0.4581

2. Probability of HeartDisease given evidence= cp

heartdisease	phi(heartdisease)
heartdisease(0)	0.3610
heartdisease(1)	0.2159
heartdisease(2)	0.1373
heartdisease(3)	0.1537
heartdisease(4)	0.1321

Result:

Thus the program was successfully executed.

Ex. No	Apply EM algorithm to cluster a set of data stored in a .CSV file. Use the same data set for clustering using the k-Means algorithm. Compare the results of these two algorithms.	Date
7		

Aim:

To apply the Expectation-Maximization (EM) algorithm and the k-Means algorithm to cluster a set of data stored in a CSV file, and then compare the results of these two algorithms.

Algorithm:

Load Dataset: Load the data from a CSV file.

Preprocess Data: Standardize or normalize the data if necessary.

Apply k-Means Algorithm: Initialize (k) cluster centroids. Assign each data point to the nearest centroid. Update centroids based on the mean of assigned points. Repeat until convergence.

Apply EM Algorithm (GMM): Initialize parameters (means, covariances, and mixture weights). E-step: Calculate the responsibilities for each data point.

M-step: Update parameters based on the responsibilities. Repeat until convergence.

Results: Use metrics such as silhouette score, inertia, or log-likelihood to compare the results.

Program:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets.samples_generator import make_blobs
X, y_true = make_blobs(n_samples=100, centers =
4,Cluster_std=0.60,random_state=0)
X = X[:, ::-1]
```

#flip axes for better plotting

```
from sklearn.mixture import GaussianMixture
gmm = GaussianMixture (n_components = 4).fit(X)
labels = gmm.predict(X)
plt.scatter(X[:, 0], X[:, 1], c=labels, s=40, cmap="viridis");
probs = gmm.predict_proba(X)
print(probs[:5].round(3))
```

```
size = 50 * probs.max(1)**2 # square emphasizes differences
plt.scatter(X[:, 0], X[:, 1], c=labels, cmap="viridis", s=size);

from matplotlib.patches import Ellipse
def draw_ellipse(position, covariance, ax=None, **kwargs):
    """Draw an ellipse with a given position and covariance"""
    Ax = ax or plt.gca()
    # Convert covariance to principal axes
    if covariance.shape == (2,2):
        U, s, Vt = np.linalg.svd(covariance)
        Angle = np.degrees(np.arctan2(U[1, 0], U[0,0]))
        Width, height = 2 * np.sqrt(s)
    else:
        angle = 0
        width, height = 2 * np.sqrt(covariance)
    #Draw the Ellipse
    for nsig in range(1,4):
        ax.add_patch(Ellipse(position, nsig * width, nsig *height,
                             angle, **kwargs))
    def plot_gmm(gmm, X, label=True, ax=None):
        ax = ax or plt.gca()
        labels = gmm.fit(X).predict(X)
        if label:
            ax.scatter(X[:, 0], X[:, 1], c=labels, s=40, cmap="viridis", zorder=2)
        else:
            ax.scatter(X[:, 0], X[:, 1], s=40, zorder=2)
        ax.axis("equal")
        w_factor = 0.2 / gmm.weights_.max()
        for pos, covar, w in zip(gmm.means_, gmm.covariances_, gmm.weights_):
```

```
draw_ellipse(pos, covar, alpha=w * w_factor)

gmm = GaussianMixture(n_components=4, random_state=42)
plot_gmm(gmm, X)

gmm = GaussianMixture(n_components=4, covariance_type="full",
random_state=42)

plot_gmm(gmm, X)
```

K-means

```
from sklearn.cluster import KMeans
```

```
#from sklearn import metrics
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
data=pd.read_csv("kmeansdata.csv")
df1=pd.DataFrame(data)
print(df1)
f1 = df1['Distance_Feature'].values
f2 = df1['Speeding_Feature'].values
```

```
X=np.matrix(list(zip(f1,f2)))
```

```
plt.plot()
plt.xlim([0, 100])
plt.ylim([0, 50])
plt.title('Dataset')
plt.ylabel('speeding_feature')
plt.xlabel('Distance_Feature')
plt.scatter(f1,f2)
plt.show()
```

```
# create new plot and data
plt.plot()
```

```

colors = ['b', 'g', 'r']
markers = ['o', 'v', 's']

# KMeans algorithm
#K = 3
kmeans_model = KMeans(n_clusters=3).fit(X)

plt.plot()
for i, l in enumerate(kmeans_model.labels_):
    plt.plot(f1[i], f2[i], color=colors[l], marker=markers[l], ls='None')
plt.xlim([0, 100])
plt.ylim([0, 50])
plt.show()

```

Driver_ID,Distance_Feature,Speeding_Feature

3423311935,71.24,28

3423313212,52.53,25

3423313724,64.54,27

3423311373,55.69,22

3423310999,54.58,25

3423313857,41.91,10

3423312432,58.64,20

3423311434,52.02,8

3423311328,31.25,34

3423312488,44.31,19

3423311254,49.35,40

3423312943,58.07,45

3423312536,44.22,22

3423311542,55.73,19

3423312176,46.63,43

3423314176,52.97,32

3423314202,46.25,35

3423311346,51.55,27

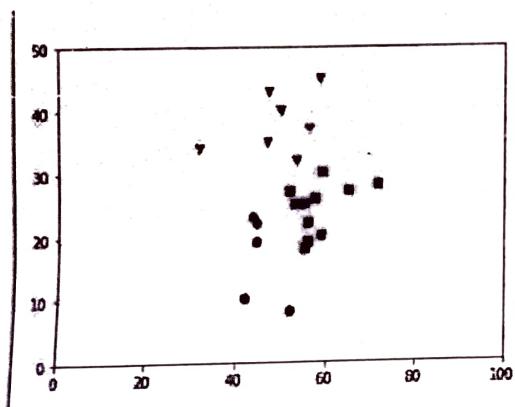
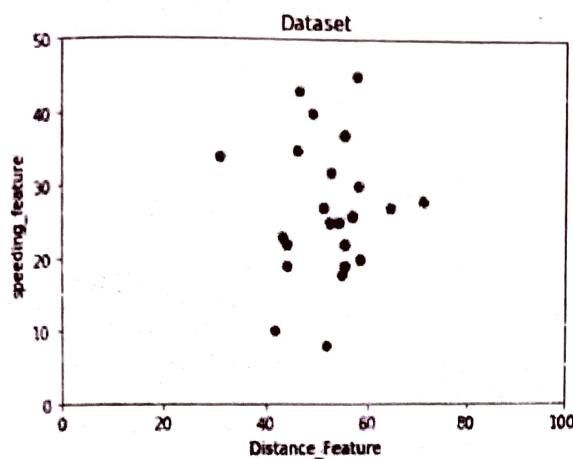
3423310666,57.05,26

3423313527,58.45,30

3423312182,43.42,23

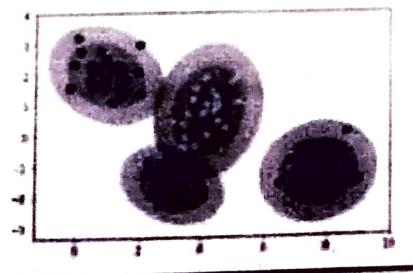
3423313590,55.68,37

3423312268,55.15,18



Output

```
[[1,0,0,0]
 [0,0,1,0]
 [1,0,0,0]
 [1,0,0,0]
 [1,0,0,0]]
```



Result:

Thus the program was successfully executed.

Ex. No	Write a program to implement k-Nearest Neighbour algorithm to classify the iris data set. Print both correct and wrong predictions.	Date
8		

Aim:

To implement the k-Nearest Neighbors (k-NN) algorithm to classify the Iris dataset and print both correct and wrong predictions. The k-NN algorithm is a simple, instance-based learning method used for classification and regression tasks.

Algorithm:

Load and Preprocess Data: Load the dataset and preprocess it (e.g., normalize features).

Distance Calculation: Compute the distance between the test instance and all training instances.

Find Neighbors: Identify the k-nearest neighbors to the test instance.

Majority Voting: Assign the class label based on the majority class among the k-nearest neighbors.

Classification: Repeat the process for all test instances to classify them.

Evaluation: Compare the predicted labels with the actual labels and print both correct and incorrect predictions.

Program:

```

import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans
import sklearn.metrics as sm
import pandas as pd
import numpy as np

iris = datasets.load_iris()

X = pd.DataFrame(iris.data)
X.columns = ['Sepal_Length','Sepal_Width','Petal_Length','Petal_Width']

```

```

y = pd.DataFrame(iris.target)
y.columns = ['Targets']
model = KMeans(n_clusters=3)
model.fit(X)
plt.figure(figsize=(14,7))
colormap = np.array(['red', 'lime', 'black'])

# Plot the Original Classifications
plt.subplot(1, 2, 1)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[y.Targets], s=40)
plt.title('Real Classification')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')

# Plot the Models Classifications
plt.subplot(1, 2, 2)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[model.labels_], s=40)
plt.title('K Mean Classification')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
print('The accuracy score of K-Mean: ',sm.accuracy_score(y, model.labels_))
print('The Confusion matrixof K-Mean: ',sm.confusion_matrix(y, model.labels_))

from sklearn import preprocessing
scaler = preprocessing.StandardScaler()
scaler.fit(X)
xsa = scaler.transform(X)
xs = pd.DataFrame(xsa, columns = X.columns)
#xs.sample(5)
from sklearn.mixture import GaussianMixture

```

```
gmm = GaussianMixture(n_components=3)
gmm.fit(xs)
y_gmm = gmm.predict(xs)

#y_cluster_gmm
plt.subplot(2, 2, 3)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[y_gmm], s=40)
plt.title('GMM Classification')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')

print('The accuracy score of EM: ',sm.accuracy_score(y, y_gmm))
print('The Confusion matrix of EM: ',sm.confusion_matrix(y, y_gmm))
```

Result:

Thus the program was successfully executed.